

ZEWEN LONG

✉Email: zamoniam2000@gmail.com • 🌐Homepage: <https://avafish.github.io> • [Google Scholar](#)

RESEARCH INTERESTS

My research interests include how to align LLMs with human preferences and values, with a focus on Recommender Systems and AI Safety.

RESEARCH EXPERIENCE

When Safety Alignment Fails to Generalize: Probing with Language Game Jailbreaks [1]

Supervisor: Prof. Shu Wu (CASIA) and Prof. Kai Chen (IIE, CAS)

- Summary: We proposed a novel jailbreak attack method to exploit large language models (LLMs) by playing custom-designed language games. This method circumvents LLM safety alignments, showcasing the vulnerability of current safety protocols.
- Contribution: Conceptualized the research idea, designed and conducted part of the experiments, authored the entire paper.
- Outcome: In submission to *The 64th Annual Meeting of the Association for Computational Linguistics (ACL 2026)*.

Graph of Thoughts Signal Modeling for Sequential Recommendation [2]

Supervisor: Prof. Shu Wu (CASIA)

- Summary: We proposed the GOT4Rec model, which first utilizes the graph of thoughts (GoT) prompting strategy in the sequential recommendation domain to capture short-term interests, long-term interests and collaborative information contained within user history sequences.
- Contribution: Designed the model, conducted all the experiments and authored the entire paper.
- Outcome: Accepted by the *2026 IEEE International Conference on Acoustics, Speech, and Signal Processing (ICASSP 2026)*.

Personalized Interest Sustainability Modeling for Sequential POI Recommendation [3]

Supervisor: Prof. Shu Wu (CASIA)

- Summary: We proposed a personalized interest sustainability model for sequential POI recommendation to capture whether each user's interest in specific POIs will sustain beyond the training time, aiding in more accurate and sustainable recommendations.
- Contribution: Designed the model, conducted all the experiments and authored part of the paper.
- Outcome: Accepted by *The 32nd ACM International Conference on Information and Knowledge Management (CIKM 2023)*.

PUBLICATIONS

- [1] Z. Long*, Y. Peng*, F. Dong, C. Li, S. Wu, and K. Chen, "When safety alignment fails to generalize: Probing with language game jailbreaks," in *Submitted to ACL Rolling Review - January 2026*, 2026.
- [2] Z. Long, L. Wang, S. Wu, Q. Liu, and L. Wang, "Got4rec: Graph of thoughts for sequential recommendation," *arXiv preprint arXiv:2411.14922*, 2024.
- [3] Z. Long*, L. Wang*, Q. Liu, and S. Wu, "Personalized interest sustainability modeling for sequential poi recommendation," in *Proceedings of the 32nd ACM International Conference on Information and Knowledge Management*, 2023, pp. 4145–4149.

WORK EXPERIENCE

Algorithm Engineer

November 2025-Present

Didi Chuxing, Beijing

- Responsible for the design and optimization of popup recommendation models deployed in the production environment.

Research Intern

July 2025-November 2025

Institute of Automation, Chinese Academy of Sciences

- Continued developing and refining my previous papers, including extended experiments and methodological improvements.
- Assisted junior lab members with research guidance, experiment design, and paper writing.

EDUCATION

Institute of Automation, Chinese Academy of Sciences (CASIA)

June 2025

M.S. in Computer Application Technology

University of Chinese Academy of Sciences (UCAS)

June 2022

B.E. in School of Cyber Security